



RD-JEPA — Recurrent Deliberation JEPA

Lens-paradigm latent world model · shared refinement F_θ reused K times · constant VRAM in K · early-exit when $V_\psi(h_k) < \tau$



Inside the Lens — on step k



$$h_k = h_{k-1} + \tanh(M_k \odot \text{MLP}_{\text{add}}(h_{k-1}, a_t)) \quad M_k = \text{Gate}(\text{MLP}_{\text{mask}}(h_k^{\text{add}})) \quad \text{VRAM constant in } K \text{ via weight sharing + gradient checkpointing}$$