





RD-JEPA v2 — Recurrent Deliberation JEPA

Kubric MOVi-A passive video · spatial latent + divergence-projection lens · curriculum $K_{\min} \rightarrow K_{\max}$ · early-exit when $V_\psi(h_k) < \tau$ · no action modality



Inside the Lens — one refinement step $F_\theta(h_{k-1}) \rightarrow h_k$ (no action input)



$$h_{\text{proj}} = h_{\text{add}} - \alpha \cdot \text{div}(h_{\text{add}}), \text{ rescaled so } \|h_{\text{proj}}\|_2 \approx \|h_{\text{add}}\|_2 \text{ (mass-preserving) } \quad \text{VRAM constant in } K \text{ via weight sharing + gradient checkpointing} \quad \text{decoder \& probe never backprop into } E_\theta \text{ or } F_\theta$$